

On the Computation of Large Spatial Datasets With the M function

Eric Marcon^{1*}

Florence Puech²

Abstract

Increasing access to large individual and spatial datasets, coupled with the development of computing power, has encouraged the search for suitable statistical tools to best analyse such data. In a recent article, Tidu et al. (2024) highlight the qualities of the M function (Marcon and Puech, 2010), a measure of spatial concentration in continuous space. They also express reservations about the computation times required. Our methodological work seeks to specify the processing of large, spatialized data sets with M using R software. Two avenues are being explored to determine the computational performance of M . Firstly, a precise evaluation of the computational time and memory requirements for geolocalised data is carried out using the *dbmss* package in R (Marcon et al., 2015) by means of performance tests. Then, as suggested by Tidu et al. (2024), we also consider the possibility of approximating the geographical positions of the entities analysed. The extent of the deterioration in the estimate of M that this approach creates are estimated and discussed, as the gains in computation time made possible by the spatial approximation of locations. The complete R code is given for the reproducibility of the results.

Keywords

Distance-based method, M-function, Performance test, R Package *dbmss*

¹AgroParisTech, UMR AMAP, CIRAD, CNRS, INRAE, IRD, Univ Montpellier, Montpellier, France.

²Université Paris-Saclay, INRAE, AgroParisTech, Paris-Saclay Applied Economics, F-91120 Palaiseau, France.

*Corresponding author: eric.marcon@agroparistech.fr

Contents

Introduction	1
1 Data simulation	2
1.1 Drawing the points	2
1.2 Gridding the space	2
2 Computing M with the <i>dbmss</i> package	3
2.1 Necessary data	3
2.2 Point pattern	3
2.3 Distance matrix	3
3 Computational performance	3
3.1 Computing time	3
3.2 Memory	3
4 Effects of approximating the position of points	4
4.1 Case of an aggregated distribution (Matérn)	4
4.2 Case of a completely random distribution (CSR)	4
5 Discussion and conclusion	4
Appendix	6
Acknowledgements	6

Introduction

Increasing access to large individual and spatial datasets and greater computing power have encouraged the development of statistical analysis tools for processing such data in the best possible way (Baddeley et al., 2016). Empirical studies at very

fine geographical levels have thus been proposed in recent years for large datasets. Particular attention has been paid to detecting the spatial structures (attraction, repulsion, independence) of individual spatialised data using analyses that are no longer based on zoned data but on geolocalised data. This type of approach has the advantage of preserving the exact positions of the entities analysed and therefore does not erase individual specificities. Various studies have shown how important it is to use this type of methodology in many fields, including geography (Sweeney and Feser, 1998; Deurloo and De Vos, 2008; Kukuliač and Horák, 2017), economics (Arbia, 1989; Marcon and Puech, 2003), ecology (Cressie, 1993; Lentz et al., 2011), biology (Dray et al., 2021), and so on. In a recent article, Tidu et al. (2024) highlight the interest of a particular statistical measure, the M function proposed by Marcon and Puech (2010). This measure, which we will refer to as M in the remainder of the article, makes it possible to highlight spatial structures within a spatialised distribution (attraction, repulsion, independence) from a study based on the distances separating the entities analysed. However, while this measure preserves all the richness of individual geolocated data, it requires a longer calculation time than other distance-based measures, since it is both a cumulative and relative measure (see Marcon and Puech, 2017, for a literature review on the advantages and limitations of a dozen

existing distance-based measures). Tidu et al. (2024) propose to limit M calculation times by introducing a voluntary positioning error for the entities analysed. For example, in their study, industrial establishments in Sardinia (Italy) are no longer located at their exact postal address but at the centroid of their municipality. This repositioning reduces calculation times, as the number of possible distances between establishments is in fact limited to the distances separating the centroids of the municipalities. This approach is similar to that of Scholl and Brenner (2015) who proposed, for the K_d function (Duranton and Overman, 2005) which characterises spatial structures using another method, to approximate the distances between pairs of entities by grouping them into classes. The method of Scholl and Brenner (2015), implemented in the *dbmss* package (Marcon et al., 2015) for R (R Core Team, 2024) provides a considerable gain in computational performance with little loss of accuracy.

In our paper, we propose to test the effectiveness of Tidu et al. (2024)’s method. First, we show the advantages of using the *dbmss* package to estimate the M function on datasets with an order of magnitude of 100,000 points or less, and we show that the computation times become excessive beyond that, on a personal computer. We then study the effect of the geographical approximation of the locations of the entities analysed. This methodological work, based on a deliberately limited number of entity locations, enables us to quantify the extent of the deterioration in information that this approach creates. These performance tests provide a precise answer to the computational advantages and limitations of the M function as a function of the size of the datasets.

The layout of the article is as follows. The first section generates the necessary data. Large point sets (of the order of several tens of thousands of points) that are either completely random or geographically concentrated are drawn. The second section details the use of the *dbmss* package to calculate the M function and its confidence interval from a table giving the position and characteristics of the points or a matrix of distances between them. The third section measures the performance of *dbmss* as a function of the size of the set of points, in terms of computing time and memory requirements. The fourth section tests the approximation which consists of grouping them together at the centre of the cells of a grid, following the approach of Tidu et al. (2024) which positions them at the centre of the administrative units in which they are located. In the last section, we conclude and discuss the advantages and the limits of an approximation of the locations on the results as well as on the computing time.

1. Data simulation

The datasets we will consider in this article are obtained by simulation. The R code is given in the ap-




Figure 1. Random draw of a set of points where the Cases (in red) are aggregated and the Controls (in blue) are distributed completely randomly. The size of the points is proportional to their weight.

pendix, which allows perfect reproducibility of the examples treated or the development of others.

1.1 Drawing the points

A set of points is drawn by a Poisson process (whose expectation of the number of points is 5,000) in a square window of side 1. Each point is assigned a qualitative mark: ‘Case’ or ‘Control’. 95% of points are ‘Controls’. 5% are ‘Cases’, whose spatial structure is studied. The weight of the points is drawn from a gamma distribution with free shape and scale parameters.

In this example, the drawing of points is completely random (*complete spatial randomness*: CSR), i.e. there is no simulation of attraction or dispersion of points which could generate spatial concentrations of points (aggregates) or, on the contrary, spatial regularities (dispersions). Sets of aggregated points can be drawn in a Matérn (1960) process.

The Cases are shown in figure 1: the aggregates are clearly visible. The controls are distributed completely randomly.

1.2 Gridding the space

Let’s consider the simulation of the Cases obtained by the Matérn process and cut the window into a grid. It simulates the usual approximation of the position of the points of an administrative unit to the position of its centre.

The approximated position of points is shown on the map in figure 2. Each cell now contains only one point of each type, whose weight is the sum of the weights of the individual points.

The values of the M function can now be calculated from the original point set or its approximation after recentring.




Figure 2. Repositioning of points in an arbitrary grid. The absence of Cases in a cell is easily detected (single-colour blue dot), as is the strong presence of Cases in a cell (two-colour dot, but predominantly red).

2. Computing M with the *dbmss* package

2.1 Necessary data

In the *dbmss* package, the function is applied to a set of points or a distance matrix. The set of points in figure 1 is used. The distance matrix between all the pairs of its points is calculated to form the data on which the performance tests will be carried out.

2.2 Point pattern

The `Mhat()` function in the *dbmss* package is used to estimate the M function. The theoretical reference value for M is 1, as this function relates the proportion of Cases up to a distance r to that observed over the entire window. The aggregation of Cases will be highlighted by values of M greater than 1 (the relative presence of Cases is greater locally than over the whole window) and the dispersion of Cases by values less than 1. We observe (figure 3) that M detects an agglomeration of Cases, which is in line with the simulation of this type of point (the controls having a completely random location on the window). The advantage of a function based on distances is clearly visible: it allows us to detect exactly at which distance(s) the attraction phenomena occur and are the most important (for functions whose values can be compared at different radii, such as M). In addition to estimating the M function, the `Menvelope()` function can be used to calculate its global confidence interval (Duran-ton and Overman, 2005) under the null hypothesis of random point location. The result is shown in figure 3.

2.3 Distance matrix

Matrices can be used to process non-Euclidean distances (transport time, road distance, etc.) which cannot be represented by a set of points. The `Mhat()` and `Menvelope()` functions are the same, and provide the same results whatever the form of the data used here




Figure 3. Value of M as a function of distance from the reference point. The confidence interval, simulated at 95%, appears in grey and is centred on the value 1.

(point set or distance matrix).

3. Computational performance

The use of the M function to characterise the spatial structure of large sets of points may be limited by the computing time or memory required.

3.1 Computing time

Calculating the distances between all pairs of points is necessary to estimate M . The calculation time is therefore expected to increase as the square of the number of points. The calculation time required for the exact calculation is evaluated for a range of numbers of points (figure 4a).

The calculation time is related to the size of the set of points by a power law. It increases less quickly than the square of the number of points. It can be estimated very precisely ($R^2 = 0.98$) by the relation $t = t_0(n/n_0)^p$ where t is the estimated time for n points (e.g.: 5.1 seconds for 100,000 points) knowing the time t_0 for n_0 points and p is the power relation (here: 1.6).

Using a distance matrix may seem an efficient way of saving computation time, but in reality calculating distances is extremely fast and the whole process from a matrix is ultimately more time-consuming. The median execution time is equal to 12 milliseconds for estimating the M function from a set of 5,000 points or 20 milliseconds for the corresponding distance matrix.

3.2 Memory

The memory used is evaluated for the same data sizes (figure 4b). **The memory required increases linearly with the number of points and is never critical for point set sizes that can be processed in reasonable times.** This highlights Tidu et al. (2024)'s conclusion about the power and computation time required when using M on large datasets. The memory used by `Dtable` objects to calculate M from a




Figure 4. Calculation time (a) in seconds and memory required (b) in MB to estimate the M function as a function of the size of the set of points. The bars represent the ± 1 standard deviation interval.

distance matrix is much greater: it is that of a numerical matrix, of the order of 8 bytes times the number of points squared, i.e. 800 MB for 10,000 points only. As the calculation time is not reduced by this approach, its use should be reserved for non-Euclidean distances.

4. Effects of approximating the position of points

Clearly, approximating the position of the points results in a loss of information: in each grid cell, the distance between all the points is set to zero, and the distance between two points in different cells is approximated by the distance between the centroids of the two cells. We therefore expect a severe error in the estimation of M on a small scale (of the order of magnitude of the size of the cells) and an error that decreases with distance, when the relative size of the cells decreases. The effect of the location approximation is first tested on a set of aggregated points, similar to the real Tidu et al. (2024) data. Secondly, the case of an unstructured set of points is considered.

4.1 Case of an aggregated distribution (Matérn)

100 sets of aggregated points (5,000 points with 5% of Cases) are simulated. To evaluate the effect of the position approximation, the exact calculation and the calculation on the grid points are performed on each set of points.

The mean values of the estimates of M are presented in figure 5a. The size of the grid cells is equal to 0.05. All neighbours at distances less than this threshold are placed at zero distance: the estimate of the function is constant up to this threshold and small-scale aggregation is underestimated. The correlation

between the M values estimated by each method is calculated at each distance (figure 5b).

The correlation is very close to 1, and the estimated values very similar, as soon as the distance taken into account exceeds the grid cell: the approximation is not a problem if the interactions between the points are studied beyond this distance. The information on interactions at short distances, i.e. within each grid cell, is lost, or, more precisely, approximated by its value at the grid scale.

4.2 Case of a completely random distribution (CSR)

The same simulations are run with a completely random set of points. The exact calculation and the calculation on the grid points are carried out on each set of points.

The average values are shown in figure 6a. The mean value of M is equal to 1 at all distances by construction: Cases and Controls are distributed completely randomly. The approximations are relatively small in value (a few percent) but artefactual aggregation is generated at small scale. As the real value of M varies little around 1, the correlations are much weaker (figure 6b) in the absence of spatial structure than in the aggregated case.

5. Discussion and conclusion

To summarise, and based on the cases proposed in this article, two conclusions can be drawn.

Firstly, it seems that **approximation on location can be considered to save computation time, given the strong correlation observed between the values of M on exact and approximated data, but keeping a very fine**

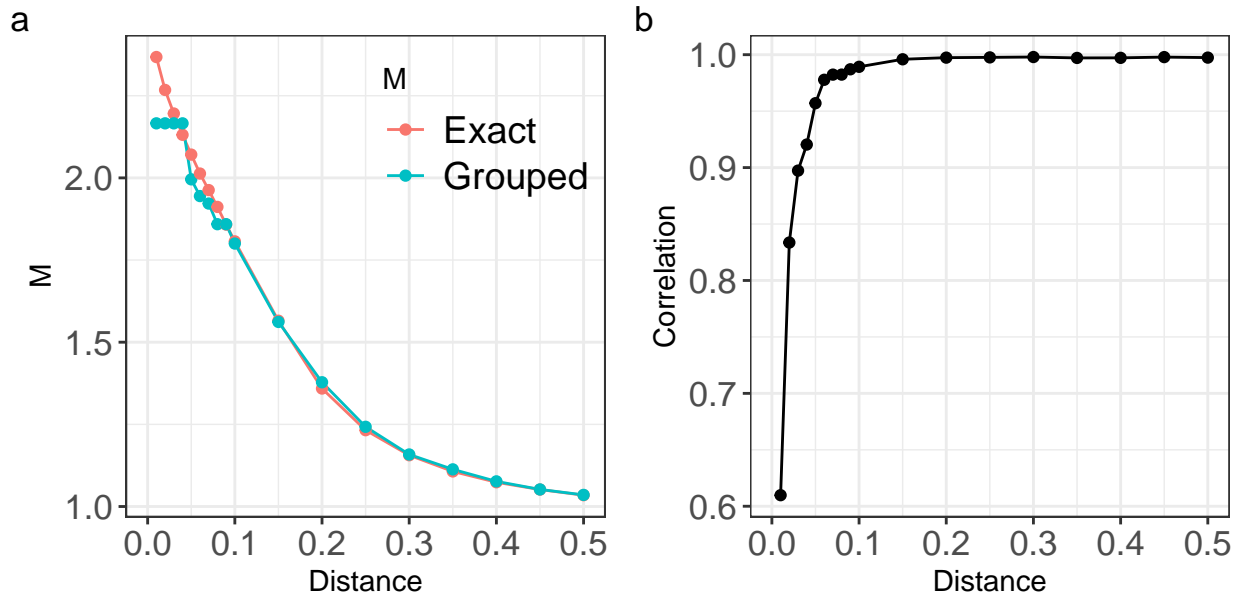


Figure 5. Average estimate of the M function from the exact position of the points compared with the values obtained by grouping the points (a) and correlation between them (b). Cases form aggregates of radius 0.1.

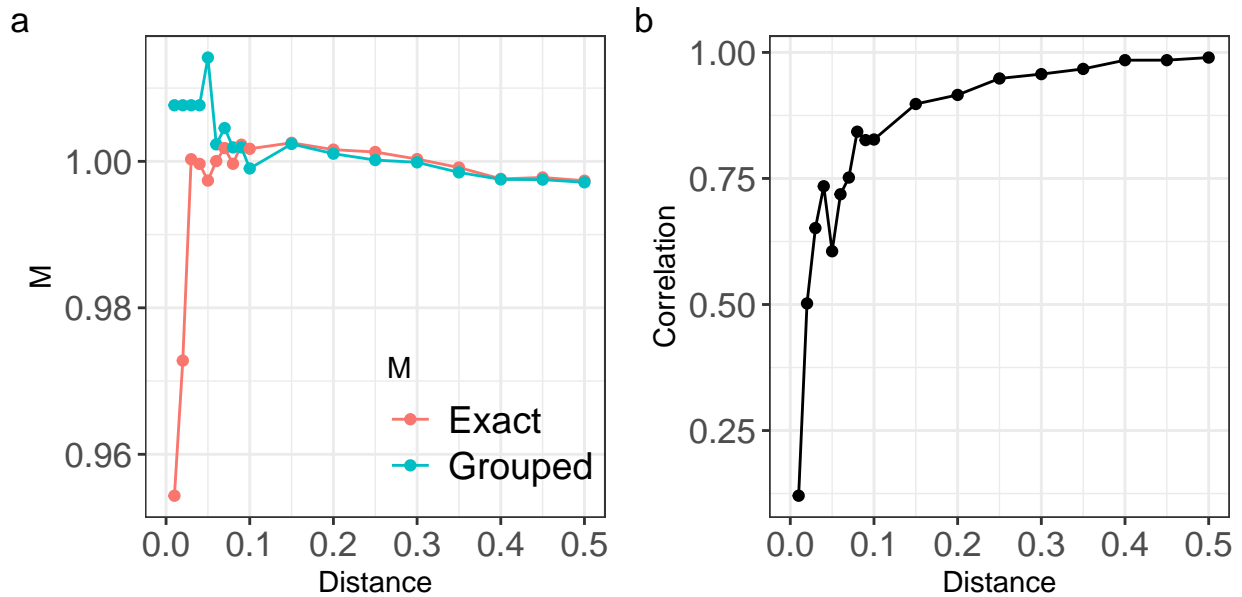


Figure 6. Average estimate of the M function from the exact position of the points compared with the values obtained by grouping the points (a) and correlation between them (b). Both Cases and Controls are drawn in a Poisson process.

grid. Our result is therefore in line with Tidu et al. (2024)’s article, which mentions strong correlations on Italian company location data. Since the spatial structure of their data is probably an intermediate case between the two cases dealt with in our article (aggregated and random theoretical distributions), the results provided by our two contributions are complementary. If the aim of the study is to look at spatial structures at very small distances, then approximating geographical positions is not desirable, because the discrepancies between the M results are greatest for these distances.

Secondly, concerning the computing burden, the calculation time for M is around 5 seconds for a set of 100,000 points on a laptop (Intel i7-1360P 2.20 GHz processor), and requires 25 MB of RAM. Calculating a confidence interval from 1,000 simulations therefore takes less than two hours. For a set of five million points, the power law predicts around an hour of computing time. 1,000 simulations would then take more than one month. Thanks to parallelization, a calculation server would drastically increase performance, but at the cost of a complexity of implementation that limits its use. If we limit ourselves to the computing power of a personal computer, exact calculation is fully justified for data of the order of 10^5 points: a few hours are enough to calculate confidence intervals. Beyond that, approximating the location reduces the size of the set of points to the number of locations selected. The price to pay is the absence of information at the scale of elementary geographical units (the grid cells in this case). **Depending on the issues addressed, a choice has to be made because this limitation may or may not be acceptable:** the overall description of the spatial structure is not significantly degraded, but the study of externalities, which is particularly interesting at short distances, is very limited.

Appendix

R code is available at the following address: <https://ericmarcon.github.io/MLargeDataSets/Appendix.pdf>

Acknowledgements

Eric Marcon benefited from an ‘Investissement d’Avenir’ grant managed by the Agence Nationale de la Recherche (LABEX CEBA, ref. ANR-10-LBX-25) and Florence Puech gratefully acknowledges financial support from INRAE.

References

- Arbia, G. (1989). *Spatial Data Configuration in Statistical Analysis of Regional Economic and Related Problems*. Dordrecht: Kluwer.
- Baddeley, A., E. Rubak, and R. Turner (2016). *Spatial Point Patterns: Methodology and Applications with*
- R*. Chapman & Hall/CRC Interdisciplinary Statistics Series. Boca Raton London New York: CRC Press.
- Cressie, N. A. (1993). *Statistics for Spatial Data*. New York: John Wiley & Sons.
- Deurloo, M. C. and S. De Vos (2008). Measuring segregation at the micro level: An application of the M measure to multi-ethnic residential neighbourhoods in Amsterdam. *Tijdschrift voor economische en sociale geografie* 99(3), 329–347.
- Dray, N., L. Mancini, U. Binshtok, F. Cheysson, W. Supatto, P. Mahou, S. Bedu, S. Ortica, E. Than-Trong, M. Krecsmarik, S. Herbert, J.-B. Masson, J.-Y. Tinevez, G. Lang, E. Beaurepaire, D. Sprinzak, and L. Bally-Cuif (2021). Dynamic spatiotemporal coordination of neural stem cell fate decisions occurs through local feedback in the adult vertebrate brain. *Cell Stem Cell* 28(8), 1457–1472.e12.
- Duranton, G. and H. G. Overman (2005). Testing for localisation using micro-geographic data. *Review of Economic Studies* 72(4), 1077–1106.
- Kukuliač, P. and J. Horák (2017). W function: A new distance-based measure of spatial distribution of economic activities. *Geographical Analysis* 49(2), 199–214.
- Lentz, J. A., J. K. Blackburn, and A. J. Curtis (2011). Evaluating Patterns of a White-Band Disease (WBD) Outbreak in *Acropora palmata* Using Spatial Analysis: A Comparison of Transect and Colony Clustering. *PLoS ONE* 6(7), e21830.
- Marcon, E. and F. Puech (2003). Evaluating the geographic concentration of industries using distance-based methods. *Journal of Economic Geography* 3(4), 409–428.
- Marcon, E. and F. Puech (2010). Measures of the geographic concentration of industries: Improving distance-based methods. *Journal of Economic Geography* 10(5), 745–762.
- Marcon, E. and F. Puech (2017). A typology of distance-based measures of spatial concentration. *Regional Science and Urban Economics* 62, 56–67.
- Marcon, E., S. Traissac, F. Puech, and G. Lang (2015). Tools to characterize point patterns: dbmss for R. *Journal of Statistical Software, Code Snippets* 67(3), 1–15.
- Matérn, B. (1960). Spatial variation. *Meddelanden från Statens Skogsforskningsinstitut* 49(5), 1–144.
- R Core Team (2024). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing.

- Scholl, T. and T. Brenner (2015). Optimizing distance-based methods for large data sets. *Journal of Geographical Systems* 17(4), 333–351.
- Sweeney, S. H. and E. J. Feser (1998). Plant size and clustering of manufacturing activity. *Geographical Analysis* 30(1), 45–64.
- Tidu, A., F. Guy, and S. Usai (2024). Measuring spatial dispersion: An experimental test on the M -index. *Geographical Analysis* 56, 384–403.